

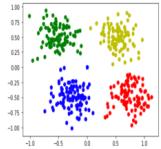
MIXTURE MODELS AND EM ALGORITHM

TOPIC 9

Name: Mirhady Dorodjatun

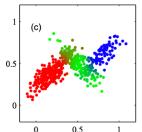
UID: u6474009

Contents



1. k-means algorithm

- Overview
- k-means optimization algorithm
- Why k-means always converge
- Choosing k and initialization method
- Python implementation



2. Gaussian Mixture

- Overview
- Latent variable modeling
- Maximum likelihood and responsibilities
- EM algorithm for Gaussian Mixtures
- Python implementation



3. Summary

Overview - k-means algorithm

- Separating dataset into k "clusters" by minimizing cost function

$$J(x, \mu) = \sum_{n=1}^N \sum_{k=1}^K r_{nk} \|x_n - \mu_k\|^2$$

↓ ↓
data point centroid

- "hard" assignment to clusters ([non-probabilistic](#))

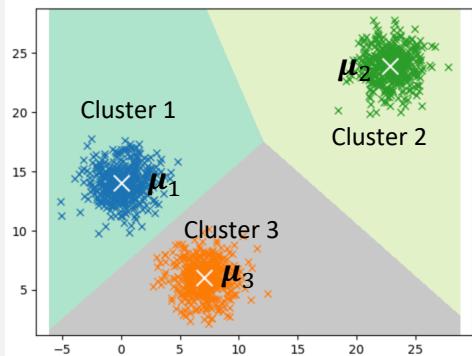


Figure 1

Goal: Minimize $L(x, \mu)$

Method: Optimization algorithm ([coordinate descent/EM](#))

- initialize value for k clusters (initialization step)
- Repeat until [convergence](#)
 - Update which cluster is assigned to each data point
 - Update the mean of the cluster

- Remember Euclidean distance represents the similarity between data points

$$\text{dist}(x, y) = \|x - y\|_2 := \sqrt{\sum_{i=1}^n (x_i - y_i)^2}$$

Overview - k-means algorithm

- Importance of k-means:
 - Classify groups of similar points without need for labeling ([unsupervised learning](#))
 - Guaranteed to converge (but may be converging to local minima) in finite steps
 - Not difficult to implement but useful: top 10 algorithm in data mining (Wu et al. 2008)
 - Can be used for image segmentation and compression
- Drawbacks:
 - Computationally expensive (Euclidean distance)
 - Time Complexity $O(n^2)$
 - Non-linearly separable dataset

k-means optimization

To optimize cost function

$$J(\boldsymbol{x}, \boldsymbol{\mu}) = \sum_{n=1}^N \sum_{k=1}^K r_{nk} \|x_n - \boldsymbol{\mu}_k\|^2$$

Steps:

1. Initialization: randomly choose k data points from \boldsymbol{x} as initial centroids $\{\boldsymbol{\mu}_1, \dots, \boldsymbol{\mu}_k\}$ (simplest one)
2. Iterate until convergence:

- a) For every $n \in \{1, \dots, N\}$, **assign to the closest centroid**

$$r_{nk} = \begin{cases} 1, & \text{if } k = \arg \min_j \|x_n - \boldsymbol{\mu}_j\|^2 \\ 0, & \text{otherwise.} \end{cases}$$

Corresponds to **E-step**
(expectation) in EM

From derivative of J to
 r_{nk} while $\boldsymbol{\mu}_k$ constant

r_{nk} determines whether x_n belongs to cluster k , and follows
"1-of- k " coding scheme

- b) For every $j \in \{1, \dots, k\}$, **update the mean (centroid) of each cluster.**

$$\boldsymbol{\mu}_k = \frac{\sum_n r_{nk} \boldsymbol{x}_n}{\sum_n r_{nk}}$$

Corresponds to **M-step**
(maximization) in EM

From derivative of J to
 $\boldsymbol{\mu}_k$ while r_{nk} constant

Illustration of k-means algorithm

E-step:

assign all data points to the closest centroid

M-step:

update the mean (centroid) of each cluster

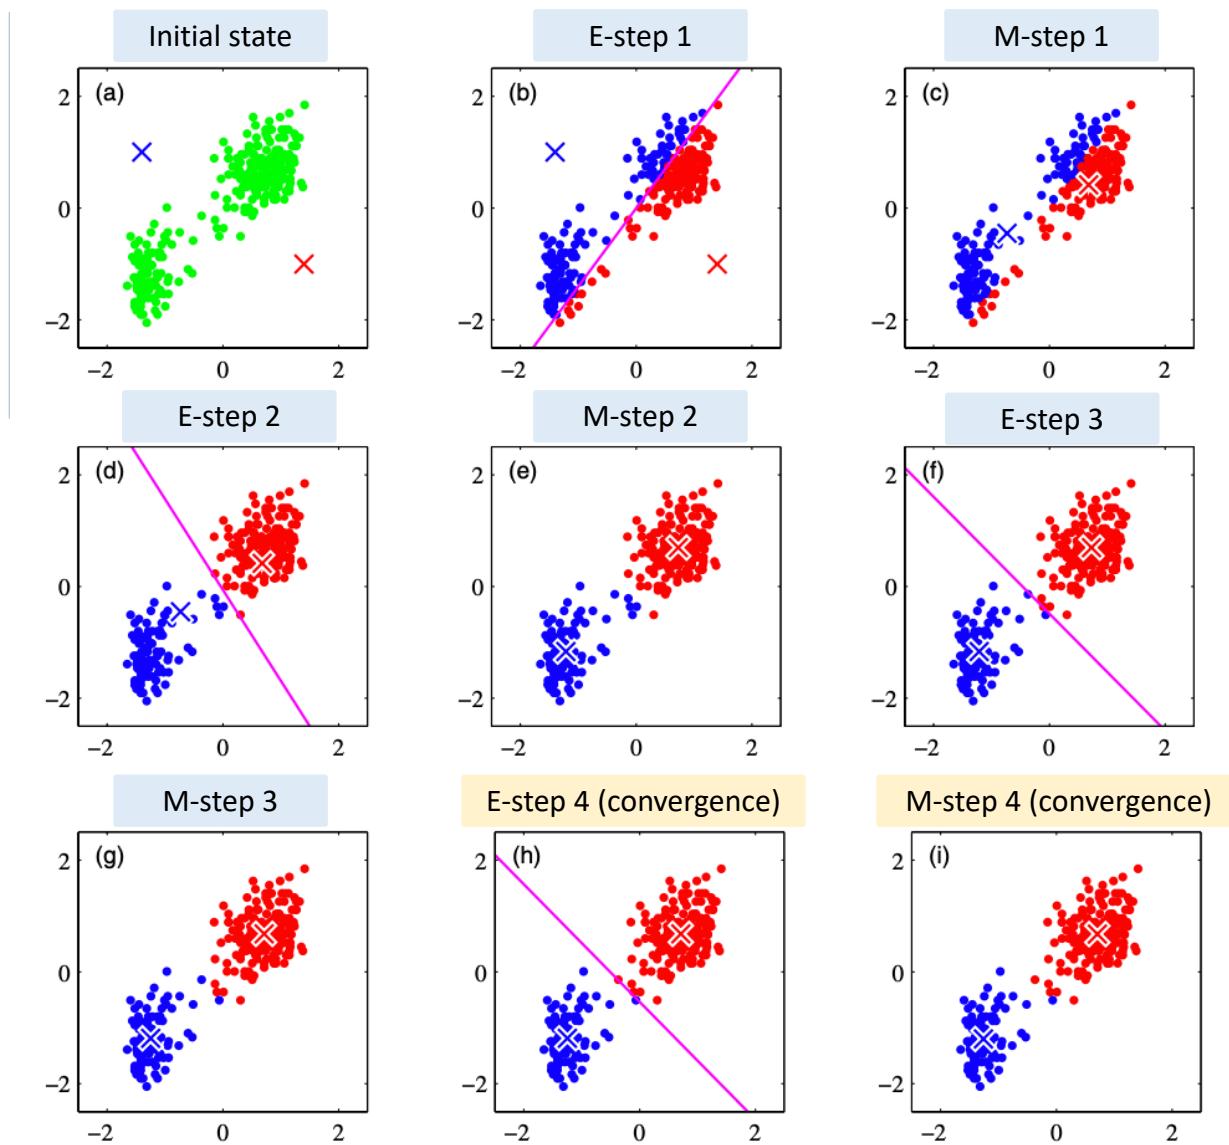


Figure 2

Cost on Convergence

- Cost (J) monotonically decrease in each EM step
- Always converge to the direction of the local minimum

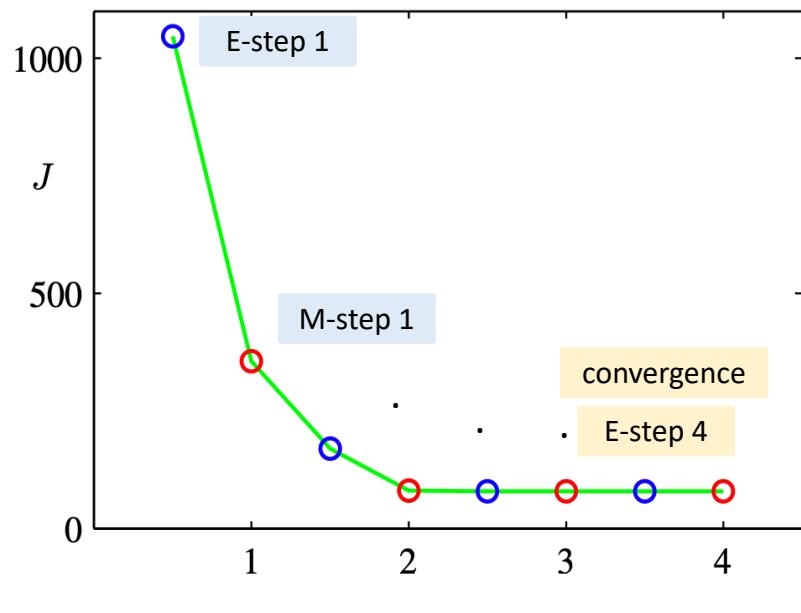


Figure 3

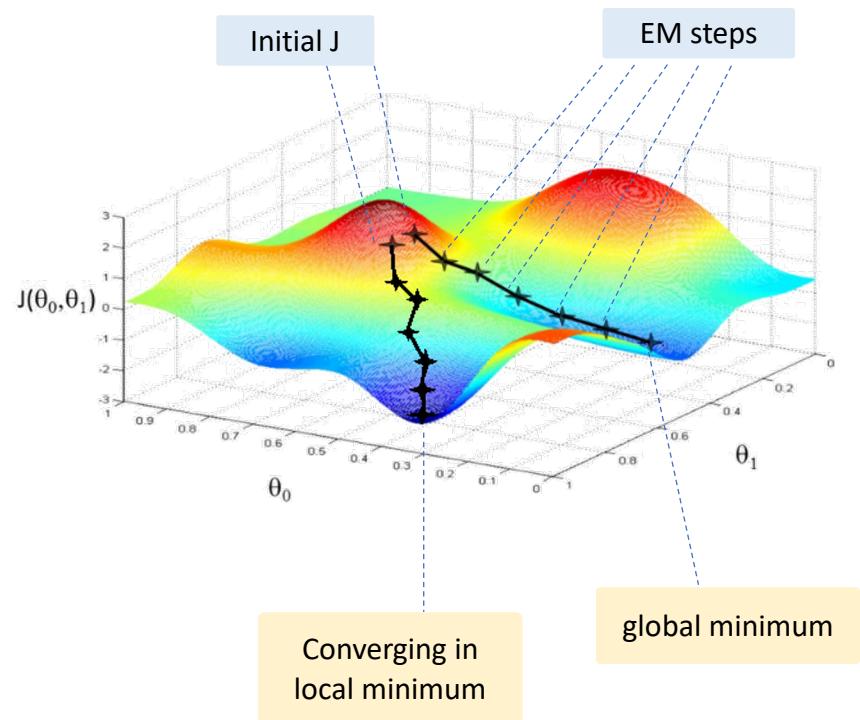


Figure 4

Why cost function monotonically decrease

Monotonically decreasing cost function will always lead to convergence.

Proof:

- E-step: assign all data points to the closest centroid

previous assignment: C_1, C_2, \dots, C_k

new assignment: C'_1, C'_2, \dots, C'_k

$$\begin{aligned} \text{cost}(C_1, C_2, \dots, C_k, z^{(1)}, \dots, z^{(k)}) &\geq \min_{C_1, \dots, C_k} \text{cost}(C_1, C_2, \dots, C_k, z^{(1)}, \dots, z^{(k)}) \\ &= \text{cost}(C'_1, C'_2, \dots, C'_k, z^{(1)}, \dots, z^{(k)}) \end{aligned}$$

- M-step: update the mean (centroid) of each cluster

previous centroid: z_1, z_2, \dots, z_k

new centroid: z'_1, z'_2, \dots, z'_k

$$\begin{aligned} \text{cost}(C'_1, C'_2, \dots, C'_k, z^{(1)}, \dots, z^{(k)}) &\geq \min_{z^{(1)}, \dots, z^{(2)}} \text{cost}(C'_1, C'_2, \dots, C'_k, z^{(1)}, \dots, z^{(k)}) \\ &= \text{cost}(C'_1, C'_2, \dots, C'_k, z'^{(1)}, \dots, z'^{(k)}) \end{aligned}$$

- General idea:

During E-step, each x_n is reassigned only when it is closer to another cluster. Therefore, E-step will either reduce or maintain the cost function (never increase it). The same holds for M-step.

Selecting value of k

- Value of k can determine result, but manually inputted

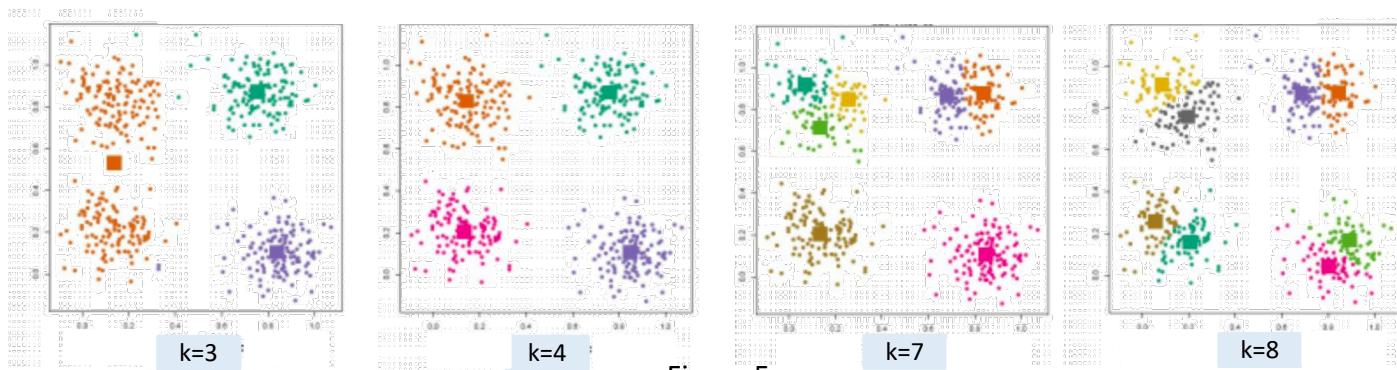


Figure 5

- Elbow method

- Observing the “elbow” from cost function result for each number of clusters.

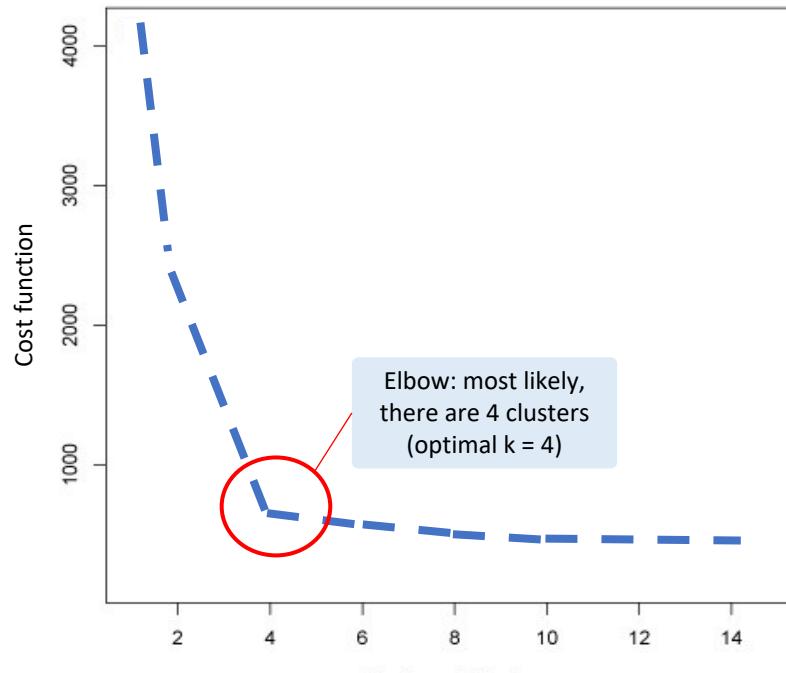


Figure 6

Importance of Initialization in k-means

- Different initialization can drastically alter the result
- Less optimal initialization method can result in empty clusters

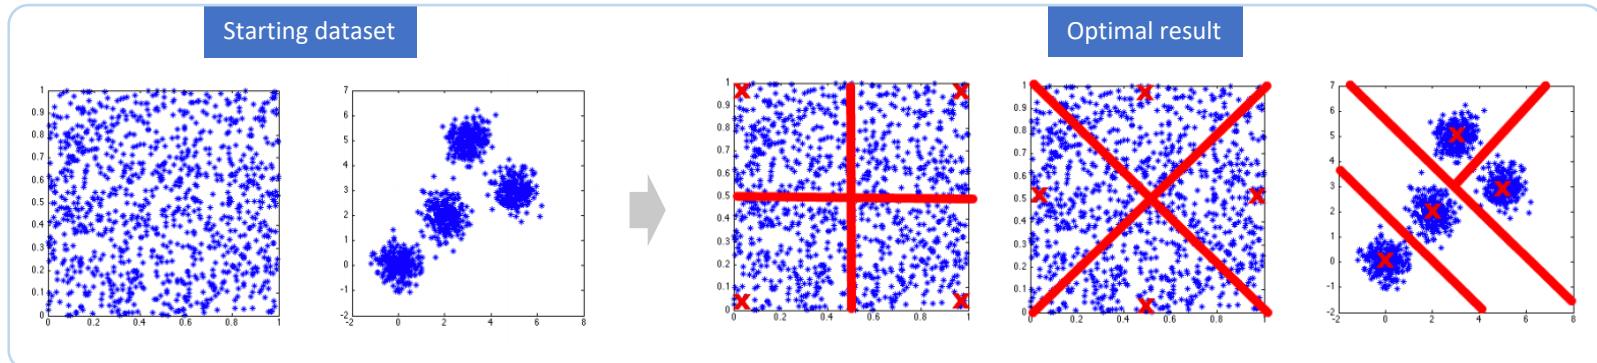


Figure 7.a

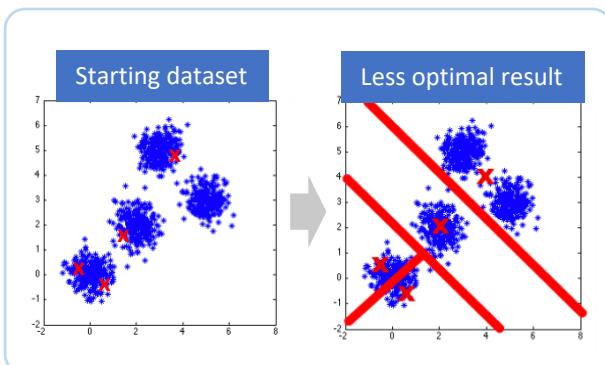


Figure 7.b

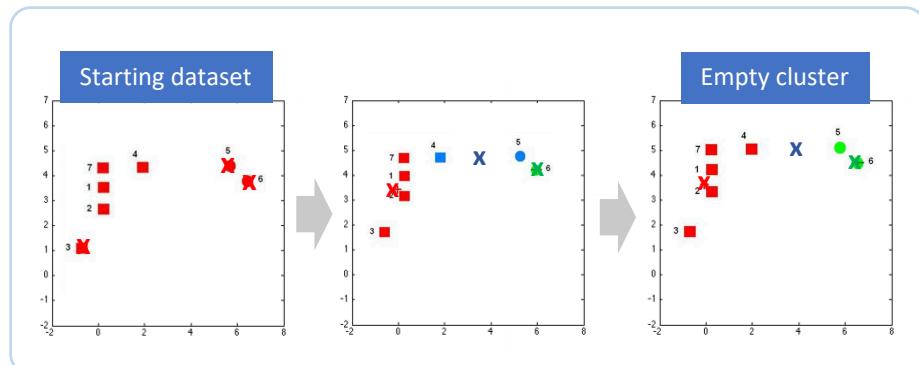
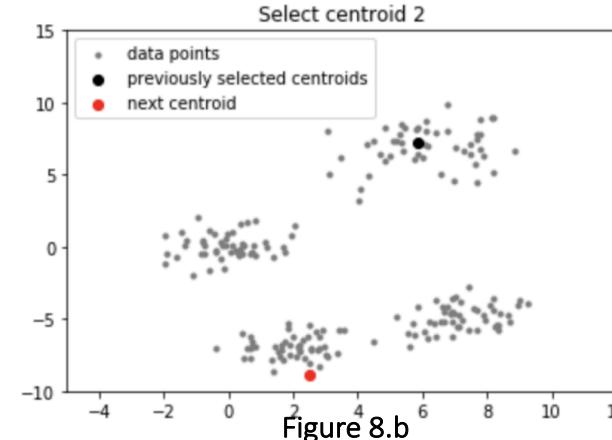
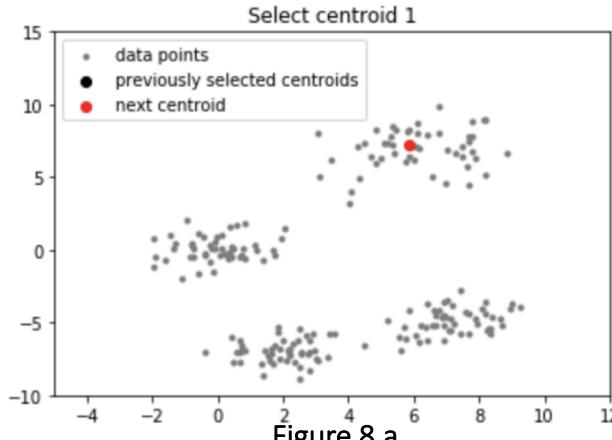


Figure 7.c

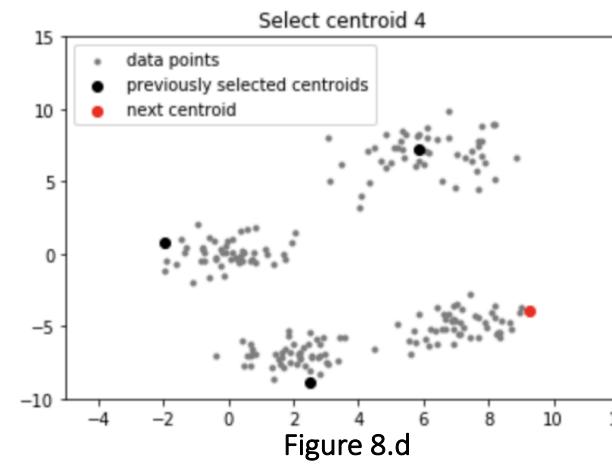
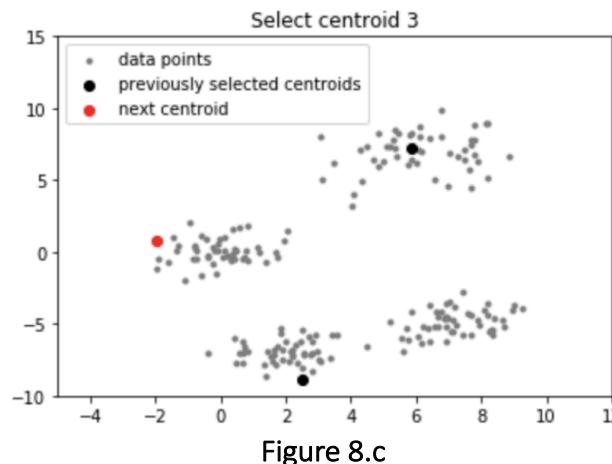
Importance of Initialization in k-means

- Example of initialization methods in k-means:

1. Random selection (simplest)
2. K-means++
 - a) Select 1 data point randomly as the first centroid
 - b) For the next centroids, choose the data point with the highest sum distance from all existing centroids
 - c) Repeat step b) until the count of centroids reaches k



The general idea is to place the centroids furthest apart with each other



K-means implementation in Python

- Initialization: random selection

```
def initialise_parameters(m, n, X):  
    C = np.zeros((m, n))  
  
    for c_index in range(0, m):  
        C[c_index] = X[np.random.randint(0,X[:,0].size)]  
  
    return C
```

- Initialization: k-means++

```
import sys  
  
def euc_distance(p1, p2):  
    return np.sum((p1 - p2)**2)  
  
def initialise_parameters(k, n, X):  
    C = []  
  
    # randomly select one data point as first centroid  
    C.append(X[np.random.randint(0,X[:,0].size)])  
  
    # then, for the remaining k-1 centroids:  
    for c_index in range(1,k):  
  
        distance = []  
        for c_index in range(X.shape[0]):  
            xn = X[c_index, :]  
            d = sys.maxsize  
  
            # compute distance of each xn from each existing centroids  
            for j in range(len(C)):  
                temp_d = euc_distance(xn, C[j])  
                d = min(d, temp_d)  
            distance.append(d)  
  
        # select data point with maximum distance as other centroids  
        distance = np.array(distance)  
        next_C = X[np.argmax(distance), :]  
        C.append(next_C)  
        distance = []  
  
    return C
```

K-means implementation in Python

- E-step

```
def E_step(C, X):
    L = np.zeros(X.shape)

    # assign all data points in X to the closest centroid
    for x_row in range(0, len(X[:,0])):
        for c_row in range(0, len(C[:,0])):
            if ( (np.linalg.norm(np.subtract(X[x_row], C[c_row])) < np.linalg.norm(np.subtract(X[x_row], L[x_row]))) or np.array_equal(L[x_row], np.zeros(L[x_row].shape)) ):
                L[x_row] = C[c_row]

    return L

L = E_step(C, X)
```

- M-step

```
def M_step(C, X, L):
    new_C = np.zeros(C.shape)

    # updates the mean (centroid) of each cluster
    for c_row in range(0, len(C[:,0])):
        n_assigned_samples = 0
        for l_row in range(0, len(X[:,0])):
            if (np.array_equal(C[c_row], L[l_row])):
                n_assigned_samples += 1
                new_C[c_row] += X[l_row]

        if (n_assigned_samples != 0):
            new_C[c_row] /= n_assigned_samples

    return new_C
```

K-means implementation in Python

- Iterating E-M steps

```
def kmeans(X, m, i):
    # calls previously defined methods
    L = np.zeros(X.shape)
    C = initialise_parameters(m, len(X[0,:]), X)

    # iterates E-step and M-step for i times
    for i in range(0, i):
        L = E_step(C, X)
        C = M_step(C, X, L)

    return C, L

C_final, L_final = kmeans(X, 4, 10)
```

- Result

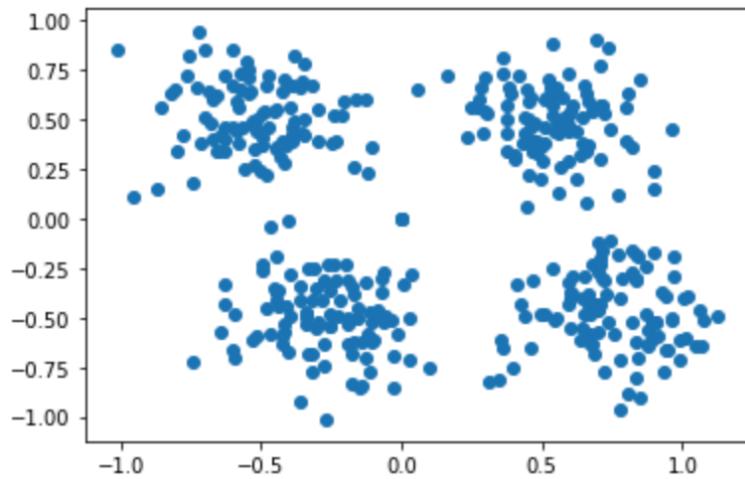


Figure 9

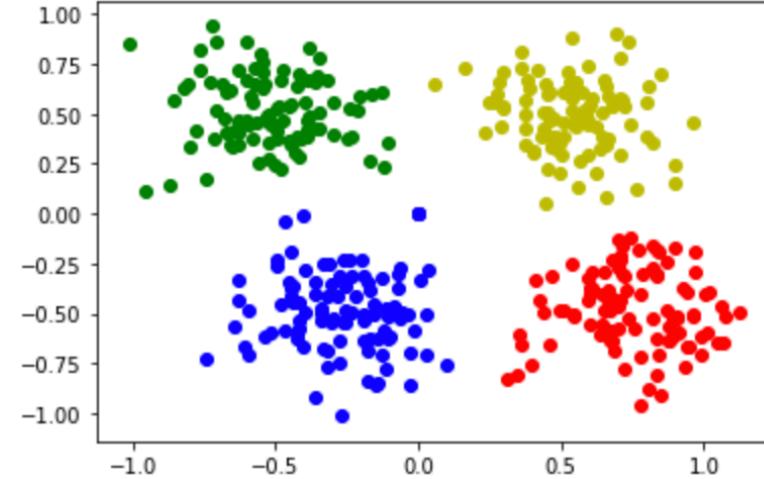


Figure 10

K-means implementation in Python

- What about data that are not linearly separable?

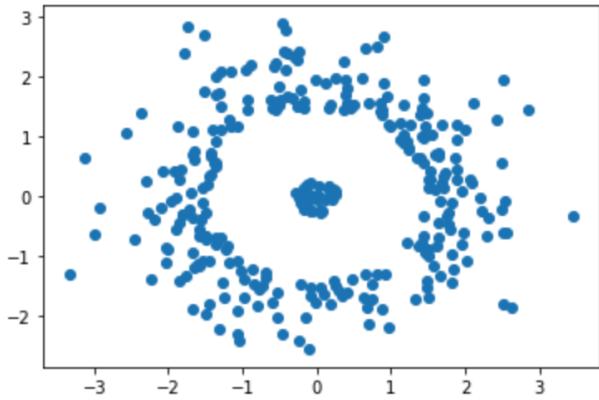
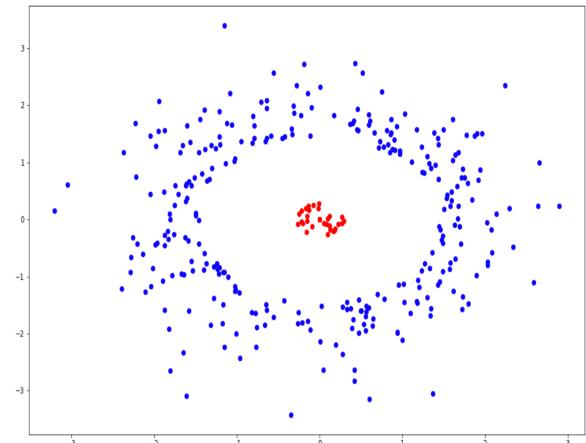


Figure 11



- Solution: add dimension to the data to make it linearly separable using k-means

```
def add_dimension(X):
    m = len(X[:,0])
    n = len(X[0,:])
    L = np.zeros(X.shape)
    X_new = np.zeros((m,n+1))

    #make X_new equals to X on every row-col of X
    for x_row in range(0, m):
        for x_col in range(0,n):
            X_new[x_row,x_col] = X[x_row,x_col]

    #make 3rd dimension equal to square of elements in X times a scalar
    for x_row in range(0, m):
        X_new[x_row, n] = (1* np.power(X_new[x_row, 0], 2))
        + (1* np.power(X_new[x_row, 1], 2))

    return X_new

Z = add_dimension(Z)
```

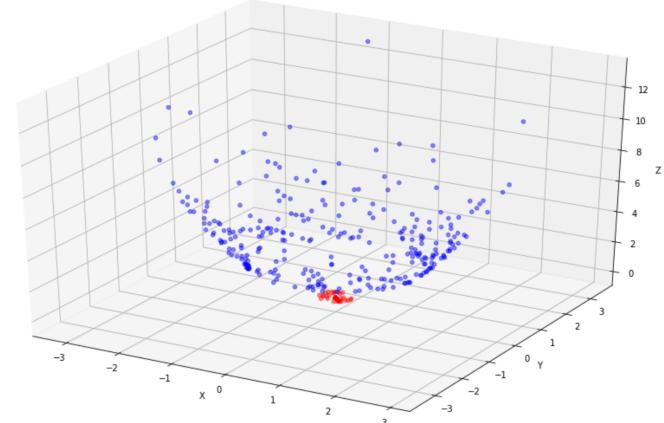


Figure 12

Image segmentation

- Segmenting image into k segments
- K-means is run on each pixels until convergence, then replace all pixel value with the centroid/mean of the cluster assigned to each pixel

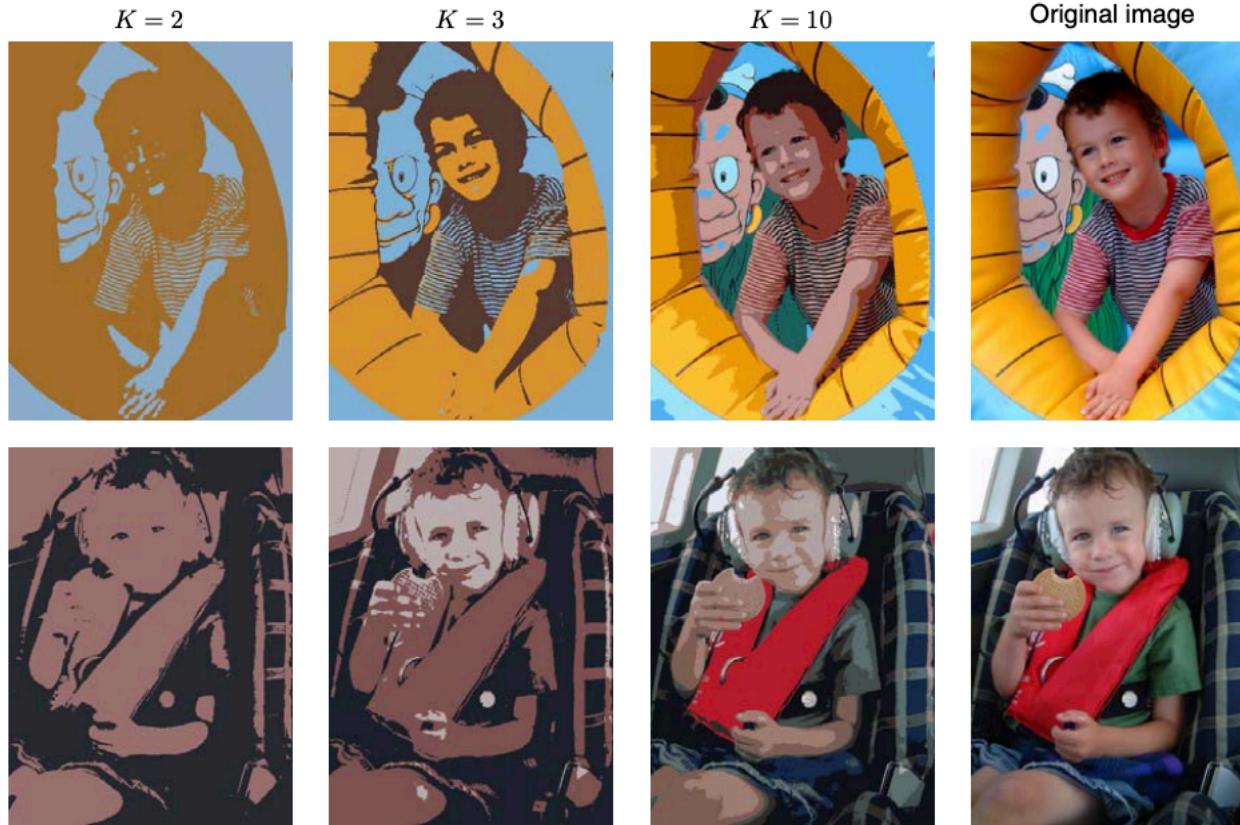


Figure 13

Image compression

- K-means as a compression algorithm ([vector quantization](#))
- Lossy compression (for reconstruction)
- K-means is run on each pixels until convergence
- Store μ_k as [code-book vectors](#), then store the labels (value of k assigned to each data point) to the code-book.

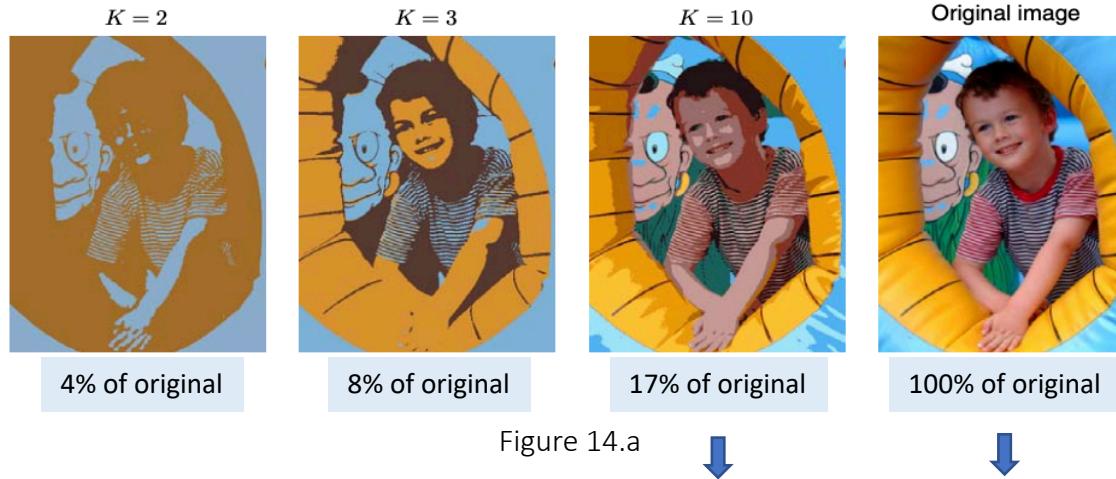


Figure 14.a

- Edge detection (using Sobel Filter):

```
im_face = cv2.imread('1.png', 0)

cv2_sobel_x = cv2.Sobel(im_face, cv2.CV_64F, 1, 0, ksize=3)
cv2_sobel_y = cv2.Sobel(im_face, cv2.CV_64F, 0, 1, ksize=3)

cv2_edge = cv2_sobel_x + cv2_sobel_y
plt.title("k=10 + sobel filter")
plt.imshow(cv2_edge, cmap='gray')
plt.show()
```

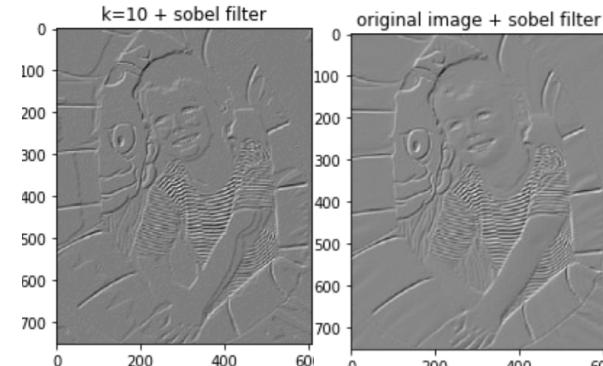
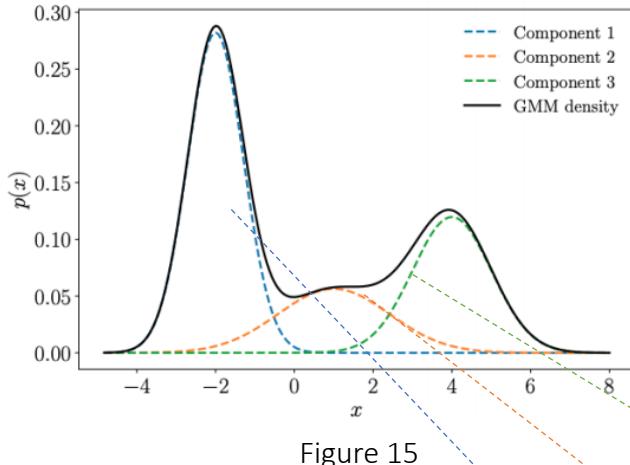


Figure 14.b

Overview – Gaussian Mixture



$$p(x) = \sum_{k=1}^K \pi_k p_k(x) \longrightarrow \text{complex to work with}$$

↓

mixture weight
(weight of the gaussian component)

introduce latent variable θ : $p(x, \theta) = p(x|\theta)p(\theta)$ \longrightarrow (product rule)

knowing $p(\theta) = \prod_{k=1}^K \pi_k^{\theta_k}$, \longrightarrow (by 1-of-k)

we then derive: $p(x|\theta) = \prod_{k=1}^K N(x|\mu_k, \Sigma_k)^{\theta_k}$

Therefore, we get the formula of Gaussian Mixture Model (GMM):

$$p(\mathbf{x}) = \sum_{\theta} p(\theta) p(\mathbf{x}|\theta) = \sum_{\theta} \prod_{k=1}^K \pi_k^{\theta_k} \prod_{k=1}^K N(\mathbf{x}|\boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)^{\theta_k} = \sum_{k=1}^K \pi_k N(\mathbf{x}|\boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)$$

in figure 15

$$p(x|\theta) = 0.5N\left(x \left| -2, \frac{1}{2}\right.\right) + 0.2N(x|1, 2) + 0.3N(x|4, 1)$$

Also, introducing the latent variable θ will allow us to **identify** (by probability) **which mixture component a data point comes from**

Overview – Gaussian Mixture

- We get the log-likelihood by: $\mathcal{L} = \log p(\mathbf{x}|\boldsymbol{\theta}) = \sum_{n=1}^N \log p(x_n|\boldsymbol{\theta}) = \sum_{n=1}^N \log \sum_{k=1}^K \pi_k \mathcal{N}(x_n|\boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)$
- Why log-likelihood instead of likelihood?
 - Log is monotonically increasing so both works equivalently for maximization
 - Adding log will help in simplifying mathematical calculations
 - Likelihood can be a very small number, which can lead to floating point error ([underflow](#))
- Importance of Gaussian Mixture
 - Modeling complex distribution with gaussian components
 - Classify groups of similar points without need for labeling ([unsupervised learning](#))
 - Richer representation of r_{nk} than k-means
- Drawbacks:
 - Even more expensive than k-means
 - Time Complexity $O(NKD^3)$ for N data points, K gaussians and D dimensions

Responsibilities

- We can get responsibilities r_{nk} using Bayes Theorem:

$$r_{nk} = p(\theta = 1|x) = \frac{p(\theta = 1)p(x|\theta = 1)}{\sum_{j=1}^K p(\theta = 1)p(x|\theta = 1)}$$

$$r_{nk} := \frac{\pi_k \mathcal{N}(x_n | \mu_k, \Sigma_k)}{\sum_{j=1}^K \pi_j \mathcal{N}(x_n | \mu_j, \Sigma_j)}$$

- Responsibility represents soft assignment in GMM, which contrasts with k-means' hard assignment:

GMM	k-means
$r_{nk} = \begin{bmatrix} 1.0 & 0.0 & 0.0 \\ 1.0 & 0.0 & 0.0 \\ 0.057 & 0.943 & 0.0 \\ 0.001 & 0.999 & 0.0 \\ 0.0 & 0.066 & 0.934 \\ 0.0 & 0.0 & 1.0 \\ 0.0 & 0.0 & 1.0 \end{bmatrix} \in \mathbb{R}^{N \times K}$	$r_{nk} = \begin{bmatrix} 1.0 & 0.0 & 0.0 \\ 1.0 & 0.0 & 0.0 \\ 0.0 & 1.0 & 0.0 \\ 0.0 & 1.0 & 0.0 \\ 0.0 & 0.0 & 1.0 \\ 0.0 & 0.0 & 1.0 \\ 0.0 & 0.0 & 1.0 \end{bmatrix} \in \mathbb{R}^{N \times K}$

- Plotted:

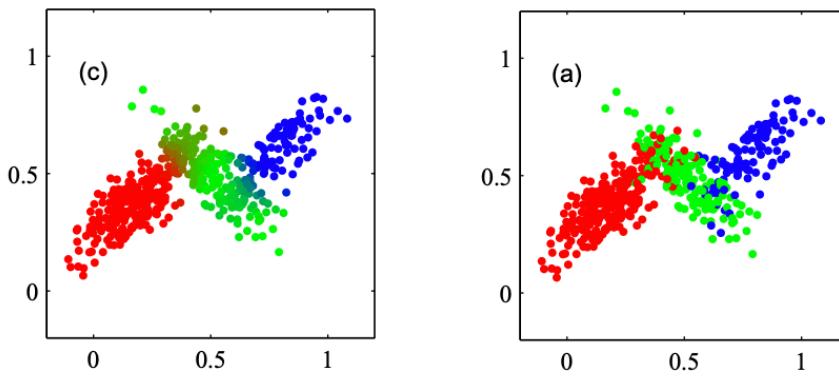


Figure 16

EM Algorithm in Gaussian Mixtures

To optimize log-likelihood function: $\mathcal{L} = \sum_{n=1}^N \log \sum_{k=1}^K \pi_k \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)$

Steps

1. Initialization:

- $\pi_k = 1/k$ for all k
- $\boldsymbol{\mu}_k$ = centroids, same as k-means algorithm
- $\boldsymbol{\Sigma}_k$ = sample variance for all k

2. Repeat until convergence:

- E-step: For every $n \in \{1, \dots, N\}$, update r_{nk} using current $\pi_k, \boldsymbol{\mu}_k$ and $\boldsymbol{\Sigma}_k$

$$r_{nk} := \frac{\pi_k \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)}{\sum_{j=1}^K \pi_j \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_j, \boldsymbol{\Sigma}_j)}$$

From Bayes Theorem In
previous slide

- M-step: For every $j \in \{1, \dots, k\}$, update parameters $\pi_k, \boldsymbol{\mu}_k$ and $\boldsymbol{\Sigma}_k$ using the updated r_{nk}

$$\begin{aligned}\boldsymbol{\mu}_k &= \frac{1}{N_k} \sum_{n=1}^N r_{nk} \mathbf{x}_n \\ \boldsymbol{\Sigma}_k &= \frac{1}{N_k} \sum_{n=1}^N r_{nk} (\mathbf{x}_n - \boldsymbol{\mu}_k)(\mathbf{x}_n - \boldsymbol{\mu}_k)^T \\ \pi_k &= \frac{N_k}{N}\end{aligned}$$

From partial derivative of log-likelihood function to $\pi_k, \boldsymbol{\mu}_k$ and $\boldsymbol{\Sigma}_k$

$$\begin{aligned}\frac{\partial \mathcal{L}}{\partial \boldsymbol{\mu}_k} = \mathbf{0}^T &\Leftrightarrow \sum_{n=1}^N \frac{\partial \log p(\mathbf{x}_n | \boldsymbol{\theta})}{\partial \boldsymbol{\mu}_k} = \mathbf{0}^T \\ \frac{\partial \mathcal{L}}{\partial \boldsymbol{\Sigma}_k} = 0 &\Leftrightarrow \sum_{n=1}^N \frac{\partial \log p(\mathbf{x}_n | \boldsymbol{\theta})}{\partial \boldsymbol{\Sigma}_k} = 0 \\ \frac{\partial \mathcal{L}}{\partial \pi_k} = 0 &\Leftrightarrow \sum_{n=1}^N \frac{\partial \log p(\mathbf{x}_n | \boldsymbol{\theta})}{\partial \pi_k} = 0\end{aligned}$$

GMM EM algorithm implementation in Python

- Initialization

```
def initialise_parameters(X, K):  
    m = X.shape[1]  
    N = X.shape[0]  
    sigma = []  
    mu = []  
    pi = []  
  
    # Calculate mean of data points X (the relative center) to use as the first centroid  
    mean_X = X.shape[0]  
    for n in range (N):  
        mean_X += X[n]  
    mean_X /= N  
  
    # Calculate stdev_X to represent the standard deviation of the datapoints X  
    stdev_X = 0  
    for n in range (N):  
        stdev_X += np.linalg.norm((X[n] - mean_X))  
    stdev_X /= N  
  
    # Now we initialize mu, first we use the mean of X as the first k  
    mu.append(mean_X)  
  
    # and then choose the next one(s) randomly, but retry if too close to mean_X  
    # we choose (2 * stdev_X) from mean_X to determine the "forbidden" range  
    for i in range(1,K):  
        isFound = False  
        while (isFound == False):  
            mu_candidate = np.asarray(X[np.random.randint(0,N),:])  
            if (np.linalg.norm((mu_candidate - mean_X)) > (2 * stdev_X)):  
                mu.append(mu_candidate)  
                isFound = True  
  
    for i in range (K):  
        # assign pi uniform for each k  
        pi.append(1/K)  
  
        # for simple positive semidefinite matrix sigma, use identity matrix  
        sigma.append(np.asarray(np.eye((m))))  
  
    return sigma, mu, pi
```

GMM EM algorithm implementation in Python

- E-step

```
def E_step(pi, mu, sigma, X):  
    N = X.shape[0]  
    m = X.shape[1]  
    K = len(pi)  
    responsibilities = np.zeros((N, K))  
    pdf = np.zeros((N,K))  
  
    # compute pdf of gaussian matrices  
    for n in range (N):  
        for k in range (K):  
            x = X[n]  
            pdf[n][k] = multivariate_normal.pdf(x, mean=np.asarray(mu[k]),  
                                              cov=sigma[k], allow_singular=True)  
  
    # compute rnk with respect to each gaussian matrix  
    for n in range (N):  
        rnk_denom = np.dot(np.asarray(pi).T, pdf[n])  
        for k in range (K):  
            responsibilities[n][k] = pi[k] * pdf[n][k] / rnk_denom  
  
    return responsibilities
```

GMM EM algorithm implementation in Python

- M-step

```
def M_step(responsibilities, X):  
    K = responsibilities.shape[1]  
    N = X.shape[0]  
    m = X.shape[1]  
    mu = []  
    sigma = []  
    pi = []  
  
    # count Nk for each gaussian distribution  
    # starting with one to prevent problems with zero value  
    # (then divided by N + 2**k later to average out the value)  
    Nk = np.ones((K,))  
    for k in range (K):  
        for n in range (N):  
            Nk[k] += responsibilities[n,k]  
    print("Nk :",Nk)  
  
    # updates mu  
    b = np.ones((responsibilities[0][0] * X[0]).shape)  
    for k in range (K):  
        for n in range (N):  
            a = np.ones((responsibilities[n][k] * X[n]).shape)  
            a = np.dot(responsibilities[n][k], X[n])  
            b += a  
    mu.append(np.asarray(b / (Nk[k])))  
  
    # updates sigma  
    c=0  
    for k in range (K):  
        for n in range (N):  
            c += responsibilities[n][k] * np.outer((X[n]) - (mu[k]), ((X[n]) - (mu[k])).T)  
    sigma.append(c / (Nk[k]))  
  
    # make all ij in sigma positive since sigma need to be positive semidefinite  
    for k in range (K):  
        for i in range (m):  
            for j in range (m):  
                if (sigma[k][i,j] < 0):  
                    sigma[k][i,j] = sigma[k][i,j] * -1  
  
    # updates pi, where we divide by N + 2**k to average out the 1 added in Nk  
    for k in range (K):  
        pi.append(Nk[k]/N + (2**k))  
  
    return mu, sigma, pi
```

GMM EM algorithm implementation in Python

- Iterating E-M steps

```
def EM(X, K, iterations):
    N = X.shape[0]

    # initialization step
    sigma, mu, pi = initialise_parameters(X, K)

    # do i iterations of EM step, by calling the predefined E_step and M_step in each iteration
    for i in range (iterations):
        print("iteration ",i, end=', ')
        responsibilities = E_step(pi, mu, sigma, X)
        mu, sigma, pi = M_step(responsibilities, X)

    return mu, sigma, pi
```

EM implementation in Python

- Result:

r_{nk} after 1 iteration

```
[[1.00000000e+000 9.43602925e-139 9.43602925e-139]
 [1.00000000e+000 6.72742056e-159 6.72742056e-159]
 [1.00000000e+000 2.29466786e-140 2.29466786e-140]
 ...
 [1.00000000e+000 5.29990521e-033 5.29990521e-033]
 [1.00000000e+000 1.65488827e-031 1.65488827e-031]
 [1.29224280e-024 5.00000000e-001 5.00000000e-001]]
```

r_{nk} after 10 iterations

```
[[2.02899408e-29 1.02437345e-02 9.89756266e-01]
 [8.70385773e-53 9.40156692e-04 9.99059843e-01]
 [9.06979976e-34 6.46505220e-03 9.93534948e-01]
 ...
 [2.99280419e-08 1.83245471e-01 8.16754499e-01]
 [1.56620844e-05 2.31746288e-01 7.68238050e-01]
 [5.79161322e-04 4.08990529e-01 5.90430310e-01]]
```

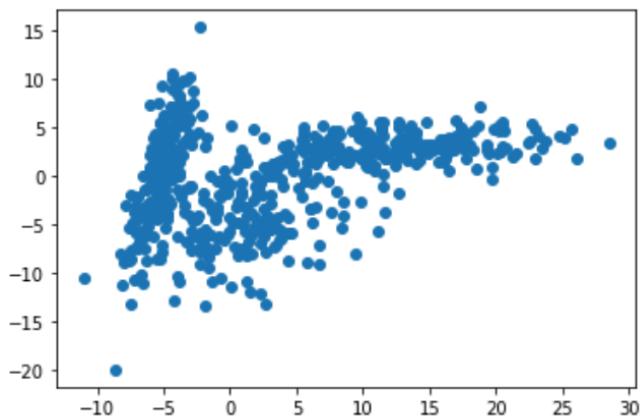


Figure 16

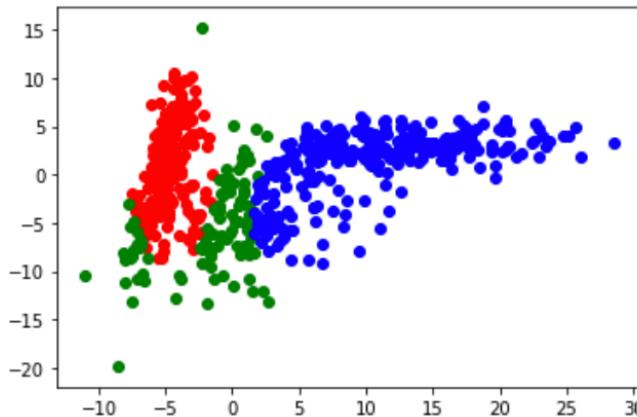


Figure 17

Limitation:
Plotting the
'intermediate'
colors

Summary

- General overview, Importance and drawback of each topic
- K-means:
 - Clustering: time complexity $O(n^2)$
 - K-means' EM algorithm and where each EM steps are derived from
 - responsibilities r_{nk} (hard assignment - nonprobabilistic)
 - Illustration of how k-means progresses with each EM steps
 - How cost function monotonically decrease but does not guarantee global minimum
 - Why the cost function monotonically decrease with each EM steps
 - Selecting value of k with elbow method
 - Importance of initialization
 - 2 initialization methods: random selection and k-means++
 - K-means for image segmentation and compression
 - Example of how the compressed image can be used, with edge detection using Sobel Filter
 - Python implementation of k-means: 2 examples, one of which with (initially) non-linearly separable dataset
- Gaussian Mixture Model:
 - Clustering: time complexity $O(NKD^3)$
 - How adding latent variable would be useful in modeling the data
 - EM algorithm and where each EM steps are derived from
 - responsibilities r_{nk} (soft assignment - probabilistic)
 - Why use log-likelihood instead of likelihood
 - Python implementation of clustering with GMM

References:

- Bishop, CM. 2006, “Pattern Recognition and Machine Learning”, New York :Springer
- Deisenroth, MP, Faisal, AA, Ong, CS. 2019, “Mathematics for Machine Learning”, Cambridge: Cambridge University Press
- COMP8600 lecture slides
- COMP6670 lecture slides

Some figures are downloaded from online sources:

- Figure 1: <https://medium.com/@saahil1292/machine-learning-103-k-means-clustering-linearly-separable-maximum-likelihood-principal-3b903bf3817c>
- Figure 8: <https://www.geeksforgeeks.org/ml-k-means-algorithm/>
- Other figures are taken from PRML book or lecture slides, except the output of python implementation

Link to ppt slide, python code implementation and video backup:

<https://www.github.com/rbdm/comp8600-video-material-topic-9>



THANK YOU

u6474009